

▼ Predicting Survival of Patient

@Author :- Priyangshu Sarkar

▼ Import Libraries

```
!pip install -q shap
```

564 kB 5.3 MB/s

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import joblib
```

```
import plotly.express as px
import plotly.offline as py
import plotly.graph_objs as go
import plotly.tools as tls
import plotly.figure_factory as ff
```

```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import (accuracy_score,
                             classification_report,
                             roc_auc_score, roc_curve, auc, precision_recall_curve,
                             confusion_matrix)
```

```
from xgboost import XGBClassifier
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import StratifiedKFold, KFold
```

```
from tqdm import tqdm
import warnings
warnings.filterwarnings('ignore')
```

```
import shap
```

```
pd.set_option('display.max_rows', 250)
```


▼ Importing Data

```
df=pd.read_csv("Dataset_Patient.csv")
```

df

	encounter_id	patient_id	hospital_id	hospital_death	age	bmi	electi
0	66154	25312	118	0	68.0	22.730000	
1	114252	59342	81	0	77.0	27.420000	
2	119783	50777	118	0	25.0	31.950000	
3	79267	46918	118	0	81.0	22.640000	
4	92056	34377	33	0	19.0	NaN	
...
45820	90508	19814	21	0	86.0	24.136910	
45821	64010	80420	21	0	55.0	27.913563	
45822	101166	124288	21	0	59.0	42.581313	
45823	27015	41654	21	0	54.0	42.850872	
45824	124157	60471	21	0	17.0	22.606103	

45825 rows × 186 columns




▼ Exploratory Data Analysis (EDA)

df.shape

(45825, 186)

df.info(verbose=True, null_counts=True)

97	h1_temp_min	38304	non-null	float64
98	d1_albumin_max	20552	non-null	float64
99	d1_albumin_min	20552	non-null	float64
100	d1_bilirubin_max	19456	non-null	float64
101	d1_bilirubin_min	19456	non-null	float64
102	d1_bun_max	40240	non-null	float64
103	d1_bun_min	40240	non-null	float64
104	d1_calcium_max	39629	non-null	float64
105	d1_calcium_min	39629	non-null	float64
106	d1_creatinine_max	40582	non-null	float64
107	d1_creatinine_min	40582	non-null	float64
108	d1_glucose_max	43288	non-null	float64
109	d1_glucose_min	43288	non-null	float64
110	d1_hco3_max	39775	non-null	float64
111	d1_hco3_min	39775	non-null	float64
112	d1_hemaglobin_max	39630	non-null	float64
113	d1_hemaglobin_min	39630	non-null	float64
114	d1_hematocrit_max	40120	non-null	float64



115	d1_hematocrit_min	40120	non-null	float64
116	d1_inr_max	17956	non-null	float64
117	d1_inr_min	17956	non-null	float64
118	d1_lactate_max	12509	non-null	float64
119	d1_lactate_min	12509	non-null	float64
120	d1_platelets_max	38958	non-null	float64
121	d1_platelets_min	38958	non-null	float64
122	d1_potassium_max	40978	non-null	float64
123	d1_potassium_min	40978	non-null	float64
124	d1_sodium_max	40371	non-null	float64
125	d1_sodium_min	40371	non-null	float64
126	d1_wbc_max	39325	non-null	float64
127	d1_wbc_min	39325	non-null	float64
128	h1_albumin_max	3671	non-null	float64
129	h1_albumin_min	3671	non-null	float64
130	h1_bilirubin_max	3451	non-null	float64
131	h1_bilirubin_min	3451	non-null	float64
132	h1_bun_max	8593	non-null	float64
133	h1_bun_min	8593	non-null	float64
134	h1_calcium_max	8149	non-null	float64
135	h1_calcium_min	8149	non-null	float64
136	h1_creatinine_max	8625	non-null	float64
137	h1_creatinine_min	8625	non-null	float64
138	h1_glucose_max	22652	non-null	float64
139	h1_glucose_min	22652	non-null	float64
140	h1_hco3_max	8449	non-null	float64
141	h1_hco3_min	8449	non-null	float64
142	h1_hemaglobin_max	9955	non-null	float64
143	h1_hemaglobin_min	9955	non-null	float64
144	h1_hematocrit_max	9783	non-null	float64
145	h1_hematocrit_min	9783	non-null	float64
146	h1_inr_max	17956	non-null	float64
147	h1_inr_min	17956	non-null	float64
148	h1_lactate_max	3994	non-null	float64
149	h1_lactate_min	3994	non-null	float64
150	h1_platelets_max	8510	non-null	float64
151	h1_platelets_min	8510	non-null	float64
152	h1_potassium_max	10183	non-null	float64
153	h1_potassium_min	10183	non-null	float64
154	h1_sodium_max	9756	non-null	float64
155	h1_sodium_min	9756	non-null	float64

df.describe()

	encounter_id	patient_id	hospital_id	hospital_death	age	
count	45825.000000	45825.000000	45825.000000	45825.000000	43665.000000	4335
mean	65625.932875	65557.129624	106.995985	0.086219	62.651804	2
std	37774.188036	37832.562156	49.579404	0.280691	16.610883	

```
df.isnull().sum(axis=0).sort_values(ascending=False)
```

```

h1_potassium_min      35642
h1_potassium_max      35642
paco2_apache          34074
paco2_for_ph_apache   34074
pao2_apache           34074
ph_apache             34074
fio2_apache           34074
d1_lactate_max        33316
d1_lactate_min        33316
d1_pao2fio2ratio_min  31707
d1_pao2fio2ratio_max  31707

d1_diasbp_invasive_max 31596
d1_diasbp_invasive_min 31596
d1_sysbp_invasive_max  31580
d1_sysbp_invasive_min  31580
d1_mbp_invasive_max    31438
d1_mbp_invasive_min    31438
bilirubin_apache       28404
d1_arterial_po2_min    28366
d1_arterial_po2_max    28366
d1_arterial_pco2_max   28346
d1_arterial_pco2_min   28346
d1_arterial_ph_max     28336
d1_arterial_ph_min     28336
h1_inr_min             27869
h1_inr_max             27869
d1_inr_min             27869
d1_inr_max             27869
albumin_apache         27571
d1_bilirubin_min       26369
d1_bilirubin_max       26369
d1_albumin_max         25273
d1_albumin_min         25273
h1_glucose_min         23173
h1_glucose_max         23173
urineoutput_apache     21481
wbc_apache             10403
bun_apache             9537
hematocrit_apache      9374
creatinine_apache      9179
sodium_apache          9161
h1_temp_min            7521
h1_temp_max            7521
hospital_admit_source   6954
d1_platelets_min       6867
d1_platelets_max       6867
d1_wbc_min             6500
d1_wbc_max             6500
d1_calcium_min         6196

```

d1_calcium_max	6196
d1_hemaglobin_max	6195
d1_hemaglobin_min	6195
d1_hco3_min	6050
d1_hco3_max	6050
d1_hematocrit_max	5705
d1_hematocrit_min	5705
d1_bun_min	5585
d1_bun_max	5585

columns that can be dropped: 'encounter_id', 'hospital_admit_source', 'icu_admit_source', 'icu_id', 'icu_stay_type', 'patient_id', 'hospital_id', 'readmission_status'

Notice the amount of missing values in each row We decide a threshold value to delete some of the attributes from the dataset (25k in this approach) **74 columns** will get deleted in this turn.

```
print("Number of rows with missing values:", df.isnull().any(axis=1).sum())
```

```
Number of rows with missing values: 45803
```

```
large_missing = df.isnull().sum(axis=0).sort_values(ascending=False)[df.isnull().sum(axis=
```

```
print("\nTotal features with more than", 25000, "missing values:", len(large_missing))
```

```
df.drop(large_missing.index.tolist() + ['encounter_id', 'hospital_admit_source', 'icu_admi
axis=1,
inplace = True)
```

```
df
```

Total features with more than 25000 missing values: 71

	hospital_death	age	bmi	elective_surgery	ethnicity	gender	height
0	0	68.0	22.730000	0	Caucasian	M	180.3
1	0	77.0	27.420000	0	Caucasian	F	160.0

df.nunique()

hospital_death	2
age	74
bmi	23169
elective_surgery	2
ethnicity	6
gender	2
height	352
icu_type	8
pre_icu_los_days	7181
weight	2566
apache_2_diagnosis	44
apache_3j_diagnosis	387
apache_post_operative	2
arf_apache	2
bun_apache	127
creatinine_apache	1037
gcs_eyes_apache	4
gcs_motor_apache	6
gcs_unable_apache	2
gcs_verbal_apache	5
glucose_apache	548
heart_rate_apache	149
hematocrit_apache	353
intubated_apache	2
map_apache	161
resprate_apache	57
sodium_apache	119
temp_apache	124
urineoutput_apache	16703
ventilated_apache	2
wbc_apache	2154
d1_diasbp_max	120
d1_diasbp_min	78
d1_diasbp_noninvasive_max	120
d1_diasbp_noninvasive_min	78
d1_heartrate_max	120
d1_heartrate_min	150
d1_mbp_max	125
d1_mbp_min	91
d1_mbp_noninvasive_max	122
d1_mbp_noninvasive_min	91
d1_resprate_max	79
d1_resprate_min	49
d1_spo2_max	31
d1_spo2_min	101
d1_sysbp_max	143
d1_sysbp_min	120
d1_sysbp_noninvasive_max	143
d1_sysbp_noninvasive_min	120

d1_temp_max	115
d1_temp_min	115
h1_diasbp_max	107
h1_diasbp_min	92
h1_diasbp_noninvasive_max	108
h1_diasbp_noninvasive_min	93
h1_heartrate_max	119
h1_heartrate_min	109
h1_mba_max	117

Removing missing values from some of the inter-related columns (bmi, weight and height)

We cut down almost 3000 instances with this process safely

```
df = df[df[['bmi', 'weight', 'height']].isna().sum(axis=1) == 0]
df
```

	hospital_death	age	bmi	elective_surgery	ethnicity	gender	height
0	0	68.0	22.730000	0	Caucasian	M	180.3
1	0	77.0	27.420000	0	Caucasian	F	160.0
2	0	25.0	31.950000	0	Caucasian	F	172.7
3	0	81.0	22.640000	1	Caucasian	F	165.1
5	0	67.0	27.560000	0	Caucasian	M	190.5
...
45819	0	87.0	35.133404	0	Caucasian	F	152.4
45820	0	86.0	24.136910	0	Hispanic	M	167.6
45821	0	55.0	27.913563	1	Caucasian	F	162.6
45822	0	59.0	42.581313	0	Caucasian	M	172.7
45823	0	54.0	42.850872	0	Caucasian	M	180.3

43355 rows × 107 columns




▼ Univariate-Multivariate Analysis

Individual plots seldom doesn't help in large datasets, in this approach let's look at the variation of instances according to each context of the column

The death rate for Male-Female patients is shown below. While the rate of **female deaths** are higher the youngest person to pass away during one of the case was a **Male of 16yrs** of age

```
fig = px.histogram(df[['age','gender','hospital_death','bmi']].dropna(), x="age", y="hospital_death",
                  marginal="box", # or violin, rug
                  hover_data=df[['age','gender','hospital_death','bmi']].columns)
fig.show()
```



▼ Average hospital death probability of patients

based on age and gender

```
age_death_F=df[df['gender']=='F'][['age','hospital_death']].groupby('age').mean().reset_index()
age_death_M=df[df['gender']=='M'][['age','hospital_death']].groupby('age').mean().reset_index()
from plotly.subplots import make_subplots
fig = make_subplots()
fig.add_trace(
    go.Scatter(x=age_death_F['age'], y=age_death_F['hospital_death'], name="Female patient")
fig.add_trace(
    go.Scatter(x=age_death_M['age'], y=age_death_M['hospital_death'], name="Male patients")
fig.update_layout()
```



```

title_text="<b>Average hospital death probability of patients<b>")
fig.update_xaxes(title_text="<b>patient age<b>")
fig.update_yaxes(title_text="<b>Average Hospital Death</b>", secondary_y=False)
fig.show()

```

Average hospital death probability of patients



▼ impacts of BMI and weight over patients


```

weight_df=df[['weight','hospital_death','bmi']]
weight_df['weight']=weight_df['weight'].round(0)
weight_df['bmi']=weight_df['bmi'].round(0)
weight_death=weight_df[['weight','hospital_death']].groupby('weight').mean().reset_index()
bmi_death=weight_df[['bmi','hospital_death']].groupby('bmi').mean().reset_index()
fig = make_subplots(rows=1, cols=2, shared_yaxes=True)
fig.add_trace(
    go.Scatter(x=weight_death['weight'], y=weight_death['hospital_death'], name="Weight"),
    row=1, col=1
)
fig.add_trace(
    go.Scatter(x=bmi_death['bmi'], y=bmi_death['hospital_death'], name="BMI"),
    row=1, col=2
)
fig.update_layout(
    title_text="<b>impacts of BMI and weight over patients<b>"
)

```

```
fig.update_yaxes(title_text="Average Hospital Death")
fig.show()
```

impacts of BMI and weight over patients



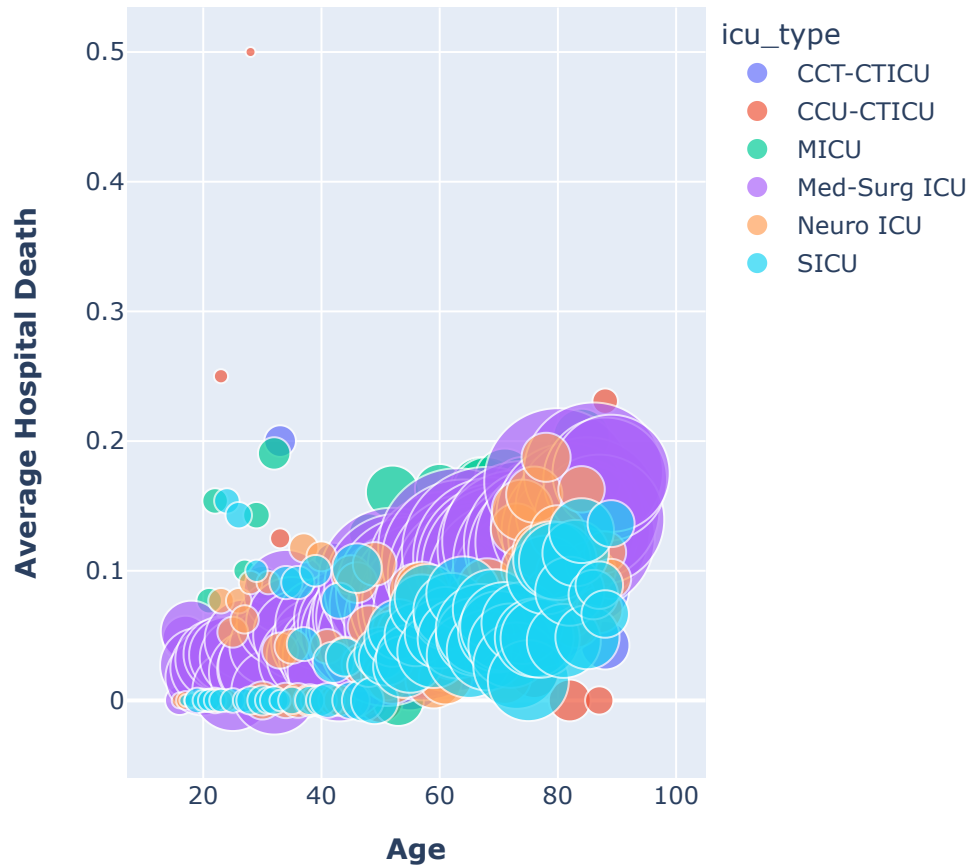
Survival rate at different types of ICU

some of the ICU wards have higher death probability pertaining to being surgical wards

```
ICU_type=df[['icu_type','age','hospital_death']]
ICU_type['icu_type']=ICU_type['icu_type'].replace({'CTICU':'CCU-CTICU',
                                                  'Cardiac ICU':'CCT-CTICU',
                                                  'CTICU':'CCT-CTICU',
                                                  'CSICU':'SICU'})
#ICU_type['pre_icu_los_days']=ICU_type['pre_icu_los_days'].round(0)
ICU_df=ICU_type.groupby(['icu_type','age']).mean().reset_index()
ICU_df['count']=ICU_type.groupby(['icu_type','age']).count().reset_index()['hospital_death']

fig = px.scatter(ICU_df, x="age", y="hospital_death", size="count", color="icu_type",
                 hover_name="icu_type", log_x=False, size_max=60,)
fig.update_layout(
    title_text="Survival rate at different types of ICU"
)
fig.update_yaxes(title_text="Average Hospital Death")
fig.update_xaxes(title_text="Age")
fig.show()
```

Survival rate at different types of ICU



▼ Hospital Death Rate, by age and Medical condition

```
apache3=df[['age','apache_3j_bodysystem','hospital_death']]
apache3=apache3.groupby(['apache_3j_bodysystem','age']).agg(['size','mean']).reset_index()

apache3['size']=apache3['hospital_death']['size']
apache3['mean']=apache3['hospital_death']['mean']

apache3.drop('hospital_death',axis=1,inplace=True)

systems =list(apache3['apache_3j_bodysystem'].unique())
data = []
list_updatemenus = []
for n, s in enumerate(systems):
    visible = [False] * len(systems)
    visible[n] = True
    temp_dict = dict(label = str(s),
                     method = 'update',
                     args = [{ 'visible': visible},
                             { 'title': '<b>'+s+'<b>' }])
    list_updatemenus.append(temp_dict)
```

```

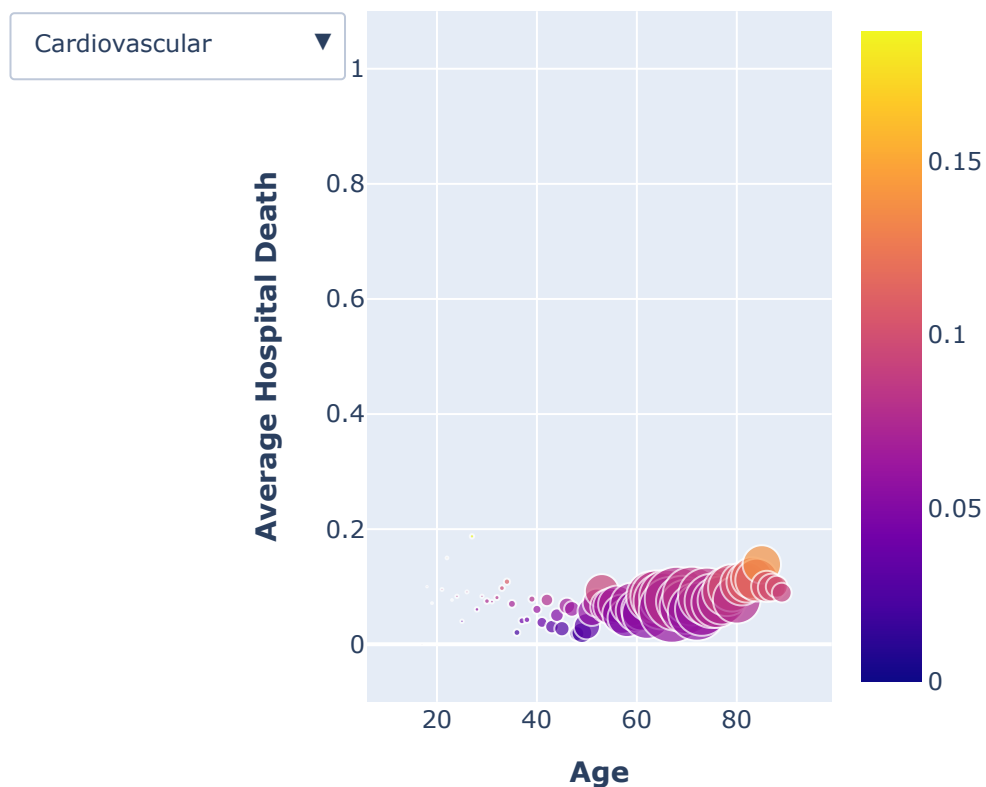
for s in systems:
    mask = (apache3['apache_3j_bodysystem'].values == s)
    trace = (dict(visible = False,
        x = apache3.loc[mask, 'age'],
        y = apache3.loc[mask, 'mean'],
        mode = 'markers',
        marker = {'size':apache3.loc[mask, 'size']/apache3.loc[mask, 'size'].sum()*1000,
            'color':apache3.loc[mask, 'mean'],
            'showscale': True}))
    )
    data.append(trace)

data[0]['visible'] = True

layout = dict(updatemenus=list([dict(buttons= list_updatemenus)]),
    xaxis=dict(title = '<b>Age<b>', range=[min(apache3.loc[:, 'age'])-10, max(ap
    yaxis=dict(title = '<b>Average Hospital Death<b>', range=[min(apache3.loc[:,
        title='<b>Survival Rate<b>' )
fig = dict(data=data, layout=layout)
py.ipplot(fig, filename='update_dropdown')

```

Survival Rate

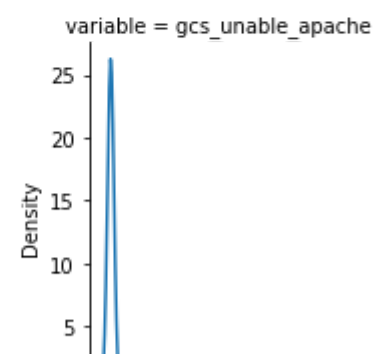
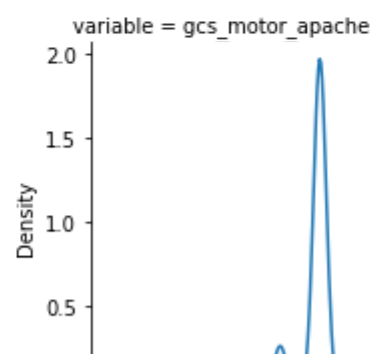
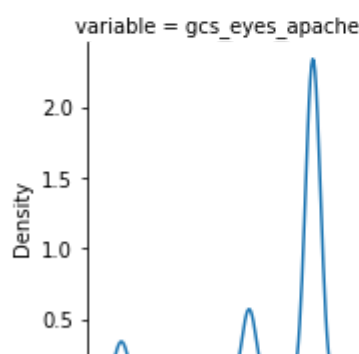
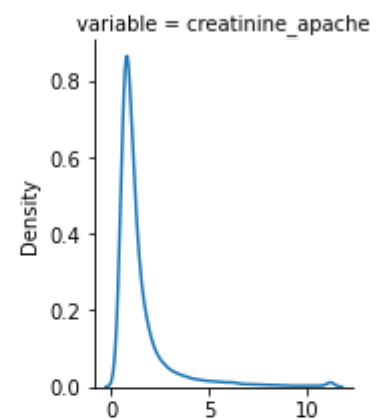


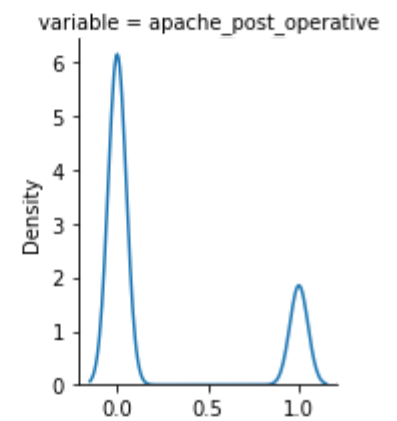
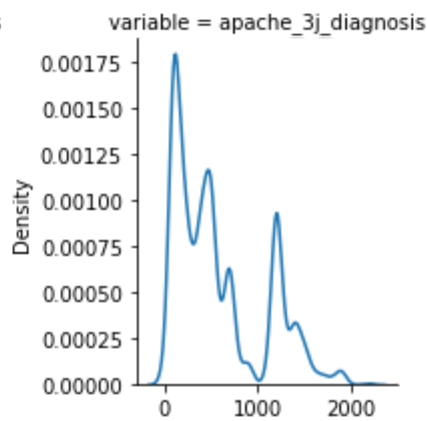
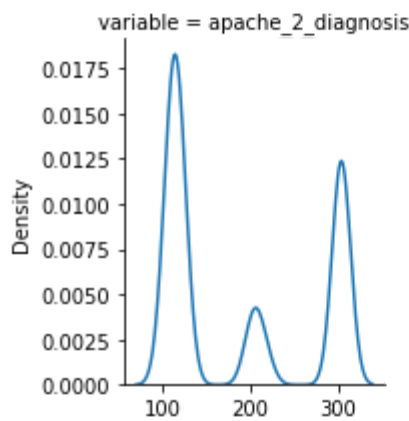


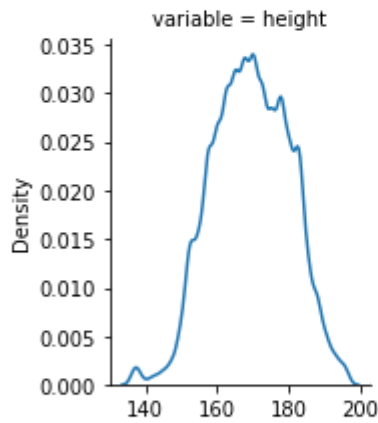





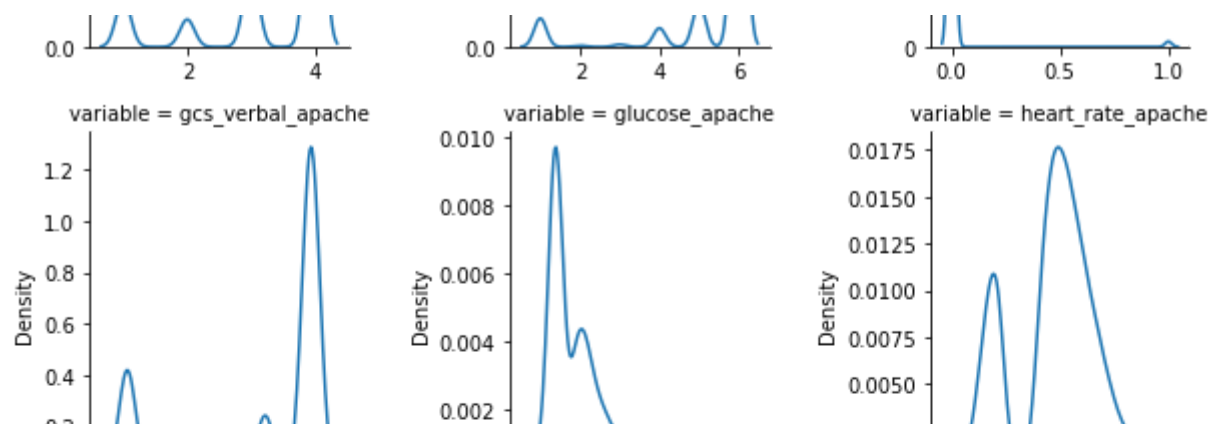
▼ Density Distribution for numerical columns

```
unpivot = pd.melt(df, df.describe().columns[0], df.describe().columns[1:])

g = sns.FacetGrid(unpivot, col="variable", col_wrap=3, sharex=False, sharey=False)
g.map(sns.kdeplot, "value")

plt.show()
```





▼ Preprocessing

converting categorical values transforming numerical columns and removing nulls

Null values for categories are replaced by **mode**, and those for numerical are replaced by **mean**.

```
numerical_cat = [  
    'elective_surgery',  
    'apache_post_operative',  
    'arf_apache',  
    'gcs_unable_apache',  
    'intubated_apache',  
    'ventilated_apache',  
    'aids',  
    'cirrhosis',  
    'diabetes_mellitus',  
    'hepatic_failure',  
    'immunosuppression',  
    'leukemia',  
    'lymphoma',  
    'solid_tumor_with_metastasis']
```

```
categorical = ['ethnicity',  
    'gender',  
    'icu_type',  
    'apache_3j_bodysystem',  
    'apache_2_bodysystem']
```

```
df.nunique()[df.nunique() == 2].index.tolist()
```

```
['hospital_death',  
    'elective_surgery',  
    'gender',  
    'apache_post_operative',  
    'arf_apache',  
    'gcs_unable_apache',  
    'intubated_apache',  
    'ventilated_apache',  
    'aids',  
    'cirrhosis',  
    'diabetes_mellitus',  
    'hepatic_failure',  
    'immunosuppression',  
    'leukemia',
```

```
'lymphoma',  
'solid_tumor_with_metastasis']
```

```
df.select_dtypes(include='O').columns.values.tolist()
```

```
['ethnicity',  
'gender',  
'icu_type',  
'apache_3j_bodysystem',  
'apache_2_bodysystem']
```

```
not_numeric = df[numerical_cat + categorical + ['hospital_death']].columns.tolist()
```

```
numeric_only = df.drop(not_numeric,axis=1).columns.tolist()
```

```
numeric_only
```

```
'd1_mbp_noninvasive_max',  
'd1_mbp_noninvasive_min',  
'd1_resprate_max',  
'd1_resprate_min',  
'd1_spo2_max',  
'd1_spo2_min',  
'd1_sysbp_max',  
'd1_sysbp_min',  
'd1_sysbp_noninvasive_max',  
'd1_sysbp_noninvasive_min',  
'd1_temp_max',  
'd1_temp_min',  
'h1_diasbp_max',  
'h1_diasbp_min',  
  
'h1_diasbp_noninvasive_max',  
'h1_diasbp_noninvasive_min',  
'h1_heartrate_max',  
'h1_heartrate_min',  
'h1_mbp_max',  
'h1_mbp_min',  
'h1_mbp_noninvasive_max',  
'h1_mbp_noninvasive_min',  
'h1_resprate_max',  
'h1_resprate_min',  
'h1_spo2_max',  
'h1_spo2_min',  
'h1_sysbp_max',  
'h1_sysbp_min',  
'h1_sysbp_noninvasive_max',  
'h1_sysbp_noninvasive_min',  
'h1_temp_max',  
'h1_temp_min',  
'd1_bun_max',  
'd1_bun_min',  
'd1_calcium_max',  
'd1_calcium_min',  
'd1_creatinine_max',  
'd1_creatinine_min',  
'd1_glucose_max',  
'd1_glucose_min',  
'd1_hco3_max',  
'd1_hco3_min',  
'd1_hemaglobin_max',  
'd1_hemaglobin_min'
```

```

d1_hemaglobin_min',
'd1_hematocrit_max',
'd1_hematocrit_min',
'd1_platelets_max',
'd1_platelets_min',
'd1_potassium_max',
'd1_potassium_min',
'd1_sodium_max',
'd1_sodium_min',
'd1_wbc_max',
'd1_wbc_min',
'h1_glucose_max',
'h1_glucose_min',
'apache_4a_hospital_death_prob',
'apache_4a_icu_death_prob']

```

```
for col in numerical_cat:
```

```
    df[col] = df[col].astype('Int64')
```

```
for col in numerical_cat:
```

```
    df[col] = df[col].fillna(df[col].mode()[0])
```

```
df[numeric_only].isna().sum(axis=0).sort_values(ascending=False)
```

```

h1_glucose_min                21686
h1_glucose_max                21686
urineoutput_apache            19517
wbc_apache                     9812
bun_apache                     8887
hematocrit_apache              8839
creatinine_apache              8572
sodium_apache                  8518
h1_temp_max                    7057
h1_temp_min                    7057
d1_platelets_min               6436
d1_platelets_max               6436
d1_wbc_max                     6077
d1_wbc_min                     6077
d1_hemaglobin_min              5800
d1_hemaglobin_max              5800
d1_calcium_max                  5717
d1_calcium_min                  5717
d1_hco3_max                     5575
d1_hco3_min                     5575
d1_hematocrit_max              5335
d1_hematocrit_min              5335
d1_bun_max                     5134
d1_bun_min                     5134
h1_mbp_noninvasive_min         5044
h1_mbp_noninvasive_max         5044
d1_sodium_max                   5007
d1_sodium_min                   5007
d1_creatinine_min               4835
d1_creatinine_max               4835
apache_4a_hospital_death_prob   4574
apache_4a_icu_death_prob        4574
d1_potassium_max                4511
d1_potassium_min                4511
glucose_apache                  4482

```

h1_diasbp_noninvasive_min	3828
h1_diasbp_noninvasive_max	3828
h1_sysbp_noninvasive_min	3820
h1_sysbp_noninvasive_max	3820
d1_glucose_min	2329
d1_glucose_max	2329
age	2045
h1_mbp_max	1862
h1_mbp_min	1862
h1_spo2_max	1656
h1_spo2_min	1656
h1_resprate_max	1562
h1_resprate_min	1562
gcs_eyes_apache	1218
gcs_motor_apache	1218
gcs_verbal_apache	1218
h1_diasbp_max	1122
h1_diasbp_min	1122
h1_sysbp_max	1115
h1_sysbp_min	1115
d1_mbp_noninvasive_max	1011
d1_mbp_noninvasive_min	1011
temp_apache	911

```
split_one = df[numeric_only].isna().sum(axis=0).sort_values()[df[numeric_only].isna().sum(
split_two = df[numeric_only].isna().sum(axis=0).sort_values()[df[numeric_only].isna().sum(
```

```
split_two
```

```
['urineoutput_apache', 'h1_glucose_max', 'h1_glucose_min']
```

```
for col in split_two:
    df[col] = df[col].fillna(df[col].mean())
```

```
process_data = df.dropna(axis=0)
```

```
process_data[categorical].nunique()
```

ethnicity	6
gender	2
icu_type	8
apache_3j_bodysystem	11
apache_2_bodysystem	10
dtype: int64	

```
icu_data = pd.get_dummies(process_data,
    prefix='isin',
    prefix_sep='_',
    columns=categorical,
    drop_first=False)
icu_data.reset_index(drop = True, inplace = True)
icu_data
```

	hospital_death	age	bmi	elective_surgery	height	pre_icu_los_days	
0	0	68.0	22.730000		0	180.30	0.541667
1	0	77.0	27.420000		0	160.00	0.927778
2	0	67.0	27.560000		0	190.50	0.000694
3	0	46.0	25.845717		0	167.60	0.000000
4	0	87.0	21.963763		0	180.30	5.046528
...
18771	0	74.0	26.096991		0	177.80	0.115278
18772	0	79.0	23.159226		0	162.56	0.059028
18773	0	85.0	18.943655		0	172.70	0.083333
18774	0	71.0	23.250046		0	177.80	0.195833
18775	0	59.0	42.581313		0	172.70	0.060417

18776 rows × 139 columns



```
icu_data.columns = [x.lower() for x in icu_data.columns.tolist()]
icu_data = icu_data.loc[:,~icu_data.columns.duplicated()]
```

```
t = icu_data['arf_apache'].dtype
for col in tqdm(icu_data.columns.tolist()):
    if icu_data[col].values.dtype == 'uint8' or t == icu_data[col].values.dtype:
        icu_data[col] = icu_data[col].astype(int)
```

100%|██████████| 133/133 [00:00<00:00, 4166.11it/s]

icu_data.dtypes

```
hospital_death      int64
age                 float64
bmi                 float64
elective_surgery     int64
height              float64
pre_icu_los_days     float64
weight              float64
apache_2_diagnosis   float64
apache_3j_diagnosis  float64
apache_post_operative int64
arf_apache           int64
bun_apache           float64
creatinine_apache    float64
gcs_eyes_apache      float64
gcs_motor_apache     float64
gcs_unable_apache    int64
gcs_verbal_apache    float64
```


glucose_apache	float64
heart_rate_apache	float64
hematocrit_apache	float64
intubated_apache	int64
map_apache	float64
resprate_apache	float64
sodium_apache	float64
temp_apache	float64
urineoutput_apache	float64
ventilated_apache	int64
wbc_apache	float64
d1_diasbp_max	float64
d1_diasbp_min	float64
d1_diasbp_noninvasive_max	float64
d1_diasbp_noninvasive_min	float64
d1_heartrate_max	float64
d1_heartrate_min	float64
d1_mbp_max	float64
d1_mbp_min	float64
d1_mbp_noninvasive_max	float64
d1_mbp_noninvasive_min	float64
d1_resprate_max	float64
d1_resprate_min	float64
d1_spo2_max	float64
d1_spo2_min	float64
d1_sysbp_max	float64
d1_sysbp_min	float64
d1_sysbp_noninvasive_max	float64
d1_sysbp_noninvasive_min	float64
d1_temp_max	float64
d1_temp_min	float64
h1_diasbp_max	float64
h1_diasbp_min	float64
h1_diasbp_noninvasive_max	float64
h1_diasbp_noninvasive_min	float64
h1_heartrate_max	float64
h1_heartrate_min	float64
h1_mbp_max	float64
h1_mbp_min	float64
h1_mbp_noninvasive_max	float64
h1_mbp_noninvasive_min	float64

▼ Modelling

```
X = icu_data.drop(['hospital_death'], axis=1)
y = icu_data['hospital_death']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30,
                                                    random_state=11,
                                                    stratify = y)
```

```
icu_data.to_csv('/content/Dataset_Patient.csv', index=False)
X_test.to_csv("/content/Dataset_Patient.csv", index=False)
```

```
y_train.value_counts()
```

```
0    11922
1     1221
Name: hospital_death, dtype: int64
```

```
y_test.value_counts()
```

```
0     5110
1      523
Name: hospital_death, dtype: int64
```

```
def modelling(X_train, y_train, X_test, y_test, **kwargs):
    scores = {}
    models = []
    if 'xgb' in kwargs.keys() and kwargs['xgb']:
        xgb = XGBClassifier()
        xgb.fit(X_train._get_numeric_data(), np.ravel(y_train, order='C'))
        y_pred = xgb.predict(X_test._get_numeric_data())
        scores['xgb'] = [accuracy_score(y_test, y_pred), roc_auc_score(y_test, y_pred)]
    #     scores['xgb']['roc_auc'] = roc_auc_score(y_test, y_pred)

    if 'rf' in kwargs.keys() and kwargs['rf']:
        rf = RandomForestClassifier(n_estimators=200)
        rf.fit(X_train, y_train)
        y_pred = rf.predict(X_test)
        scores['rf'] = [accuracy_score(y_test, y_pred), roc_auc_score(y_test, y_pred)]
    #     scores['rf']['roc_auc'] = roc_auc_score(y_test, y_pred)
        models.append(rf)

    if 'extree' in kwargs.keys() and kwargs['extree']:
        extree = ExtraTreesClassifier()
        extree.fit(X_train, y_train)
        y_pred = extree.predict(X_test)
        scores['extree'] = [accuracy_score(y_test, y_pred), roc_auc_score(y_test, y_pred)]
    #     scores['extree']['roc_auc'] = roc_auc_score(y_test, y_pred)
        models.append(extree)

    return scores
```

```
modelling(X_train,y_train, X_test, y_test, xgb=True, rf=True, extree=True)
```

```
{'extree': [0.9227303446722235, 0.6084453975078975],
 'rf': [0.9249831043027709, 0.6350888225888226],
 'xgb': [0.9247578283397162, 0.6482187341562342]}
```

```
def model_performance(model, y_test, y_hat) :
    conf_matrix = confusion_matrix(y_test, y_hat)
    trace1 = go.Heatmap(z = conf_matrix ,x = ["0 (pred)","1 (pred)"],
                        y = ["0 (true)","1 (true)"],xgap = 2, ygap = 2,
                        colorscale = 'Viridis', showscale = False)

    #Show metrics
    tp = conf_matrix[1,1]
```

```

fn = conf_matrix[1,0]
fp = conf_matrix[0,1]
tn = conf_matrix[0,0]
Accuracy = ((tp+tn)/(tp+tn+fp+fn))
Precision = (tp/(tp+fp))
Recall = (tp/(tp+fn))
F1_score = (2*(((tp/(tp+fp))*(tp/(tp+fn)))/((tp/(tp+fp))+(tp/(tp+fn)))))

show_metrics = pd.DataFrame(data=[[Accuracy , Precision, Recall, F1_score]])
show_metrics = show_metrics.T

colors = ['gold', 'lightgreen', 'lightcoral', 'lightskyblue']
trace2 = go.Bar(x = (show_metrics[0].values),
                y = ['Accuracy', 'Precision', 'Recall', 'F1_score'], text = np.round_(s
                textposition = 'auto',
                orientation = 'h', opacity = 0.8,marker=dict(
                color=colors,
                line=dict(color='#000000',width=1.5)))

#Roc curve
model_roc_auc = round(roc_auc_score(y_test, y_hat) , 3)
fpr, tpr, t = roc_curve(y_test, y_hat)
trace3 = go.Scatter(x = fpr,y = tpr,
                    name = "Roc : " + str(model_roc_auc),
                    line = dict(color = ('rgb(22, 96, 167)'),width = 2), fill='tozeroy
trace4 = go.Scatter(x = [0,1],y = [0,1],
                    line = dict(color = ('black'),width = 1.5,
                    dash = 'dot'))

# Precision-recall curve
precision, recall, thresholds = precision_recall_curve(y_test, y_hat)
trace5 = go.Scatter(x = recall, y = precision,
                    name = "Precision" + str(precision),
                    line = dict(color = ('lightcoral'),width = 2), fill='tozeroy')

#plots
model = model

#Subplots
fig = tls.make_subplots(rows=2, cols=2, print_grid=False,
                        specs=[
#                                [{'colspan': 2}, None],
#                                [{}, {}],
#                                [{}, {}],
#                                [{'colspan': 2}, None]
                                ],
                        subplot_titles=('Confusion Matrix',
                                      'Metrics',
                                      'ROC curve'+ " "+ '('+ str(model_roc_auc)+' )',
                                      'Precision - Recall curve',
                                      ))

fig.append_trace(trace1,1,1)
fig.append_trace(trace2,1,2)

```

```

fig.append_trace(trace3,2,1)
fig.append_trace(trace4,2,1)
fig.append_trace(trace5,2,2)

fig['layout'].update(showlegend = False, title = '<b>Model performance report</b><br>'
                      autosize = False, height = 1500,width = 830,
                      plot_bgcolor = 'rgba(240,240,240, 0.95)',
                      paper_bgcolor = 'rgba(240,240,240, 0.95)',
                      margin = dict(b = 195))
fig["layout"]["xaxis2"].update((dict(range=[0, 1])))
fig["layout"]["xaxis3"].update(dict(title = "false positive rate"))
fig["layout"]["yaxis3"].update(dict(title = "true positive rate"))
fig["layout"]["xaxis4"].update(dict(title = "recall"), range = [0,1.05])
fig["layout"]["yaxis4"].update(dict(title = "precision"), range = [0,1.05])
fig["layout"]["xaxis5"].update(dict(title = "Percentage contacted"))
fig["layout"]["yaxis5"].update(dict(title = "Percentage positive targeted"))
fig.layout.titlefont.size = 14

py.iplot(fig)

```

▼ Parameter Tuning

```
gkf = KFold(n_splits=3, shuffle=True, random_state=42).split(X=X_train, y=y_train)
```

```

fit_params_of_xgb = {
    "early_stopping_rounds":100,
    "eval_metric" : 'auc',
    "eval_set" : [(X_test, y_test)],
    'verbose': 100,
}

```

```

# A parameter grid for XGBoost
params = {
    'booster': ["gbtree"],
    'learning_rate': [0.1],
    'n_estimators': range(100, 500, 100),
    'min_child_weight': [1],
    'gamma': [0],
    'subsample': [0.8],
    'colsample_bytree': [0.8],
    'max_depth': [5],
    "scale_pos_weight": [1]
}

```

```

xgb_estimator = XGBClassifier(
    objective='binary:logistic',
    # silent=True,
)

```

```
gsearch = GridSearchCV(
```