

```
!nvidia-smi
```

```
NVIDIA-SMI has failed because it couldn't communicate with the NVIDIA driver. Make su
```



```
!pip install -q git+https://github.com/tensorflow/docs
```

```
Building wheel for tensorflow-docs (setup.py) ... done
```

```
import numpy as np
import pandas as pd
```

```
import tensorflow as tf
import tensorflow_hub as hub
import tensorflow_datasets as tfds
import seaborn as sns
import matplotlib.pyplot as plt
plt.rcParams['figure.figsize'] = (12, 8)
from IPython import display
```

```
import pathlib
import shutil
import tempfile
```

```
import tensorflow_docs as tfdocs
import tensorflow_docs.modeling
import tensorflow_docs.plots
```

```
print("Version: ", tf.__version__)
print("Hub version: ", hub.__version__)
print("GPU is", "available" if tf.config.list_physical_devices('GPU') else "NOT AVAILABLE")
```

```
#logdir = pathlib.Path(tempfile.mkdtemp())/ "tensorboard_logs"
#shutil.rmtree(logdir, ignore_errors=True)
```

```
☞ Version: 2.8.2
Hub version: 0.12.0
GPU is NOT AVAILABLE
```

```
df = pd.read_csv("/content/train.csv.zip",
                 compression='zip',
                 low_memory=False
)
```

```
df.shape
```

```
(1306122, 3)
```

```
df.head()
```

```
df.target.plot(kind='hist', title="Target Distribution")
```

```
from sklearn.model_selection import train_test_split
```

```
train_df, remaning = train_test_split(df, random_state=42, train_size=0.01, stratify=df.ta
```

```
train_df.head()
```

```
remaning.head()
```

```
validation_df, _ = train_test_split(remaning, random_state=42, train_size=0.001, stratify=
print(f'train size: {train_df.shape}\nvalidation size: {validation_df.shape}')
```

```
    train size: (13061, 3)
    validation size: (1293, 3)
```

```
validation_df.head()
```

```
train_df.target.head(15).values
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0])
```

```
train_df.question_text.head(15).values
```

```
array(['What is your experience living in Venezuela in the current crisis? (2018)',
      'In which state/city the price of property is highest?',
      'Do rich blacks also call poor whites, "White Trash"?',
      'Should my 5 yr old son and 2 yr old daughter spend the summer with their
father, after a domestic violent relationship?',
      'Why do we have parents?',
      'Do we experience ghost like Murphy did in Interstellar?',
      'Are Estoniano women beautiful?'],
      dtype=object)
```

```

        'There was a Funny or Die video called Sensitivity Hoedown that got pulled. Does anyone know why?',
        'Is it a good idea to go in fully mainstream classes, even if I have meltdowns that might disrupt people?',
        'What classifies a third world country as such?',
        'Is being a pilot safe?',
        'Who is Illiteratendra Modi? Why does he keep with him a Rs 1 lakh pen?',
        'Have modern management strategies such as Total supply Chain Management applied to education? Can they be?',
        'Why are Lucky Charms considered good for you?',
        'How many people in India use WhatsApp, Facebook, Twitter and Instagram?'],
        dtype=object)

```

```

module_url = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1" # @param ["https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1", "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"]

```

```

import time
import os
def get_log_path(log_dir="logs/fit"):
    uniqueName = time.strftime("log_%Y_%m_%d_%H_%M_%S")
    log_path = os.path.join(log_dir, uniqueName)
    print(f"savings logs at: {log_path}")

    return log_path

log_dir = get_log_path()

    savings logs at: logs/fit/log_2022_06_26_19_42_42

tensorboard_cb = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
early_stopping_cb = tf.keras.callbacks.EarlyStopping(monitor='val_loss', patience=2, mode='min')

EpochDots_cb = tfdocs.modeling.EpochDots()

CALLBACKS_LIST = [tensorboard_cb, early_stopping_cb, EpochDots_cb]

def train_and_evaluate_model(module_url, embed_size, name, trainable=False):
    hub_layer = hub.KerasLayer(module_url, input_shape=[], output_shape=[embed_size], dtype='float32')
    model = tf.keras.models.Sequential([
        hub_layer,
        tf.keras.layers.Dense(256, activation=tf.nn.relu),
        tf.keras.layers.Dense(64, activation=tf.nn.relu),
        tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)
    ])
    model.compile(
        optimizer = tf.keras.optimizers.Adam(learning_rate=0.0001),
        loss = tf.keras.losses.BinaryCrossentropy(),
        metrics = [tf.keras.metrics.BinaryAccuracy(name='accuracy')]
    )

    model.summary()

    history = model.fit(

```

```

train_df.question_text,
train_df.target,
epochs = 100,
validation_data = (validation_df.question_text, validation_df.target),# validation
callbacks = CALLBACKS_LIST,
verbose=0
)

```

histories = {} # will saave history object returned by train_and_evaluate_model. So I can

```

module_url = "https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1" #@param ["https://tfhub.dev/google/tf2-preview/gnews-swivel-20dim/1"]

```

```

histories['gnews-swivel-20dim'] = train_and_evaluate_model(module_url, embed_size=20, name

```

Model: "sequential"

Layer (type)	Output Shape	Param #
keras_layer (KerasLayer)	(None, 20)	400020
dense (Dense)	(None, 256)	5376
dense_1 (Dense)	(None, 64)	16448
dense_2 (Dense)	(None, 1)	65
Total params: 421,909		
Trainable params: 21,889		
Non-trainable params: 400,020		

```

Epoch: 0, accuracy:0.9313, loss:0.2716, val_accuracy:0.9381, val_loss:0.2009,
.....

```

```

module_url = "https://tfhub.dev/google/tf2-preview/nnlm-en-dim50/1" #@param ["https://tfhub.dev/google/tf2-preview/nnlm-en-dim50/1"]

```

```

histories['nnlm-en-dim50'] = train_and_evaluate_model(
    module_url, embed_name='nnlm-en-dim50', trainable=False
)

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
keras_layer_1 (KerasLayer)	(None, 50)	48190600
dense_3 (Dense)	(None, 256)	13056

dense_4 (Dense)	(None, 64)	16448
dense_5 (Dense)	(None, 1)	65

```
=====
Total params: 48,220,169
Trainable params: 29,569
Non-trainable params: 48,190,600
```

```
Epoch: 0, accuracy:0.9102, loss:0.3493, val_accuracy:0.9381, val_loss:0.2273,
.....
```

```
module_url = "https://tfhub.dev/google/tf2-preview/nnlm-en-dim128/1" # @param ["https://tfh
module_url": "https://tfhub.dev/google/
```

```
histories['nnlm-en-dim128'] = train_and_evaluate_model(module_url, embed_size=128, name='
```

```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
=====		
keras_layer_2 (KerasLayer)	(None, 128)	124642688
dense_6 (Dense)	(None, 256)	33024
dense_7 (Dense)	(None, 64)	16448
dense_8 (Dense)	(None, 1)	65

```
=====
Total params: 124,692,225
Trainable params: 49,537
Non-trainable params: 124,642,688
```

```
Epoch: 0, accuracy:0.9346, loss:0.3229, val_accuracy:0.9381, val_loss:0.2105,
.....
```

